

Homework 4

ECE345 - Group 16

November 24th, 2023

Total pages: 9

Team Member	Student Number
Ardavan Alaei Fard	1007934620
Rowan Honeywell	1007972945
Isaac Muscat	1007897135

Contents

Question 1	2
(a)	2
(b)	2
Question 2	3
Question 3	4
Question 4	5
Question 5	8
(a)	8
(b)	9

Question 1

(a)

To show that G and H share the same set of minimum spanning trees, we can show that for any sub-MST, they both share the same safe edges. In theorem 23.1 in CLRS 3rd edition, it is shown that the light edge crossing the cut of any cut of G that respects A is safe for A . The theorem uses the fact that since (u, v) is a light edge, $w(u, v) \leq w(x, y)$ for another edge $w(x, y)$ crossing the cut. This means that swapping $w(u, v)$ for $w(x, y)$ leads to $w(T') = w(T) - w(x, y) + w(u, v) \leq w(T)$. If both $w(x, y)$ and $w(u, v)$ get incremented by one when translating G to H , the above theorem shows that (u, v) is still the safe edge for G and H since the inequality remains unchanged. Therefore, G and H share the same minimum spanning trees.

(b)

G and H do not share the same shortest paths between all pairs of vertices.

As a counter example, consider a graph G with a destination vertex v_d and source vertex v_s . We need only to show that the shortest path from v_s to v_d is not the same for both graphs. Let there be two paths connecting the two vertices. In graph G , one path consists of 8 edges with weights of 1 while the other path consists of one edge with a weight of 9. In graph G , the shortest path involves traversing the path with 8 edges for a cost of 8 instead of traversing the single edge with a cost of 9. In contrast, in graph H the 8 edge path has each edge weighted at 2 with a path cost of 16 and the single edge is weighted at 10. Therefore, the shortest path will be achieved by traversing the single edge. Therefore, the shortest path from v_s to v_d is different in G and H .

Question 2

The shortest path from every node in North America v_{NA} to every node in Europe v_E , must include the transatlantic edge (e_{ta}). Additionally, both v_{NA1} and v_{E1} will be in the set of vertices for each shortest path since it is the only path connecting the two subgraphs G_{NA} and G_E .

Using the above statements, the shortest path from each $v_{NA} \in G_{NA}$ to each $v_E \in G_E$ can be broken down into the following:

$$\text{SHORTESTPATH}(v_{NA} \rightarrow v_{NA1}) + \text{SHORTESTPATH}(v_E \rightarrow v_{E1}) + e_{ta}$$

The shortest path from $v_{NA} \rightarrow v_{NA1}$ can be found in $\mathcal{O}((V + E) \lg V)$ using Dijkstra's algorithm (CLRS 3rd Ed, pg 661 - 662) with e_{ta} removed. Similarly, the shortest path from $v_E \rightarrow v_{E1}$ can be found in $\mathcal{O}((V + E) \lg V)$ using Dijkstra's algorithm with e_{ta} removed. The shortest path can then be found by concatenating the previous two paths found with e_{ta} .

Since two $\mathcal{O}((V + E) \lg V)$ operations are run at most (in fact the amortized cost is less - not asymptotically - because both subgraphs only run Dijkstra's algorithm on their nodes and edges which is less than E and V), the asymptotic time complexity of the algorithm is $\mathcal{O}((V + E) \lg V)$. Dijkstra's algorithm uses a priority-queue of vertices to determine the next vertex with the smallest cost which takes $\mathcal{O}(V)$ space.

Question 3

First, take each conversion rate and construct a directed graph using an adjacency list. Each node represents a currency and its edges represents the conversion ratio from the currency of node A to node B. For instance, inserting the conversion CAD to USD involves creating two edges with each edge specifying the conversion from and to both currencies. An adjacency list is used because it keeps BFS and DFS down to $\mathcal{O}(V + E)$ in time complexity and it takes asymptotically less space and time to create than an adjacency matrix. To resolve a currency exchange query, a DFS or BFS (both take $\mathcal{O}(V + E)$ time) is run on the first node in the exchange query and terminating if the second node is found in the search. If DFS/BFS fail to find the second currency, then there is no conversion possible. If the search is successful, the currency exchange rate will be the multiplication of each edge weight. The algorithm takes $\mathcal{O}(V + E)$ time because for adjacency lists, inserting an edge and a vertex takes constant time and $V + E$ insertions are done to construct the graph. The time complexity of the search is also $\mathcal{O}(V + E)$. The algorithm takes $\mathcal{O}(V + E)$ space because only an adjacency list is used to store the currency exchange rates and their relationships.

Question 4

In this question, we have three constraints for our flow network:

1. Each doctor should only be assigned to work on the days when they are available.
2. For a given parameter c , each doctor should be assigned to work at most c vacation days.
3. For each vacation period j , any given doctor should be scheduled at most one day in D_j .

First, we solve the problem without considering the third condition. Assume for doctor i , A_i would represent the set of vacation days that the doctor is available to work. To begin with, the source node is connected to the column of doctor nodes (d_i for doctor i). The edge for this connection is set to c_i for each doctor. This way, we ensure that each doctor would work less or equal to their capacity, covering constraint number two. Then, each doctor node is connected to the next layer which is the vacation days (v_{jm} for representing the m th day of j th vacation period), but only for the days that they are available (i.e. there is no connection if the day is not part of A_i). The edge capacity for this connection is set to 1, representing one whole day in the vacation period. The vacation days are then connected to the sink with an edge capacity of 1. A simplified version of this network flow is presented in Figure 1 with three doctors and two vacation periods. For solving the problem, we need to set a demanding value of N (the total number of vacation days) for the sink's inflow. In other words, as we solve for the maximum flow, if the inflow of the sink would be equal to N , we know that we have a feasible schedule with at least one doctor for every vacation day. If the sink value would be k units less than N , it means that there are k vacation days that

we do not have a doctor available for (i.e. there is no flow in k vacation days).




Figure 1: Flow network solution ignoring constraint number 3.

Now that we have a valid graph for the easier version, it is time to tackle constraint number 3. For this part, we have to ensure that each doctor works only on one day during a vacation period. To clarify this constraint using our previous example in Figure 1, doctor d_1 has to work either on v_{11} or v_{12} regardless of how many days they can work during the vacations (c_1). For doing this, we need a new layer to limit the flow that goes into each vacation period from each doctor node. Thus, we introduce a layer between the doctors and the vacation days to represent vacation period availability for each doctor (p_{ij} is the node for i th doctor availability for j th vacation period). The edges are again all set to one to ensure the doctor works for only one day per vacation period (i.e. since the inflow of a vacation period is one, only one of the vacation days would have a flow). The new version of our flow is represented

in Figure 2, derived from our previous example. Again for completing the algorithm, we set a demanding value of N (the total number of vacation days) for the sink's inflow, knowing that our schedule is feasible as long as this requirement (demand) is met in our circulation. As we use the Ford-Fulkerson polynomial-time algorithm for this question, it requires at most N iterations to find all the augmenting paths in the residual network. Thus, according to the textbook, the overall time complexity would be $\mathcal{O}(EN)$ (E is the total number of edges) plus the required time for checking if the sink's input flow is equal to N or not.




Figure 2: Flow network full solution, derived from Figure 1 (For more clarity, only the edges with weight not equal to 1 are specified).

Question 5

(a)

To preface, we will characterize a “bad” swap as the swapping of a heavy child to a right child position. For each left-heavy node (more descendants on its left than its right) that we come across in the traversal of the right-most branch, it is guaranteed that a bad swap will occur as it will be swapped to a right child position.

Next, because we only ever traverse rightwards when merging, every time we come across a left-heavy node, the **maximum** possible number of nodes that remain to be traversed is divided by two. Since this is the maximal case, there are **at most** $\lg n$ left-heavy nodes in the right-most branch and consequently a maximum of $\lg n - 1$ bad swaps that can occur (since we do not swap the last node).

Charge & Allocation: Let us charge $2 \lg n - 1$ credits for each merge operation, where n is the total number of nodes in both h and h' . We will **deposit** one credit for each of the first $\lg n$ nodes that are merged along the right-most branch. Additionally, we will **deposit** one credit for each bad swap that is performed. Finally, we will **charge** one credit for each swap (good or bad).

Credit Invariant: For each bad swap and corresponding heavy child in a right child location, there will be an extra credit stored. Also, one credit will be stored for each of the first $\lg n$ nodes that must be merged and subsequently swapped.

Argument for Positive Credit: A perfect **leftist** heap will have a right-most branch of maximum size $\lg n$. Since we deposited one credit for each of the first $\lg n$ merged nodes, the cost of swapping each of these nodes will be covered. Of course, there may be additional

nodes in the right-most branch since skewed heaps swap unconditionally. Each bad swap creates a longer right-sided branch somewhere in the heap which may be traversed in a future merge operation. It is important to note that the next time this branch is traversed, it will engage a **good** swap, as the right-heavy node will be swapped to the left. Since one credit was deposited for each bad swap, and there are a maximum of $\lg n - 1$ bad swaps that can occur (as shown previously), these self-correcting swaps will be covered by what was previously deposited.

Therefore, all of the swapping operations can be accounted for by an initial deposit of $\lg n + \lg n - 1 = 2\lg n - 1$ credits. It has thus been shown that SKEWEDHEAPMERGE has an amortized time-complexity of $\mathcal{O}(\lg n)$.

(b)

Analysis of INSERT: In the case of inserting a new node into a skewed heap, we can simply employ the SKEWEDHEAPMERGE algorithm from part (a). This works because inserting a single node into a skewed heap equivalent to merging two heaps, the only condition being that one is only a root node.

Analysis of DELETEMIN: Similar to the analysis of INSERT, we can also use SKEWEDHEAPMERGE to delete the minimum element from a skewed heap. This can be done by disconnecting the heap's root node (minimum) from its two children and merging the resulting left and right sub-heaps.

In conclusion, since both INSERT and DELETEMIN can be performed through a single merge, they both have the same amortized time-complexity of $\mathcal{O}(\lg n)$.